

'''

USAGE:

python cam_test.py

'''

```
import torch
import joblib
import torch.nn as nn
import numpy as np
import cv2
import argparse
import torch.nn.functional as F
import time
import cnn_models

from torchvision import models

# load label binarizer
lb = joblib.load('../outputs/lb.pkl')

model = cnn_models.CustomCNN()
# model = cnn_models.CustomCNN().cuda()

model.load_state_dict(torch.load('../outputs/model.pth'))
print(model)
print('Model loaded')

def hand_area(img):
    hand = img[100:324, 100:324]
    hand = cv2.resize(hand, (224, 224))
    return hand

cap = cv2.VideoCapture(0)

if (cap.isOpened() == False):
    print('Error while trying to open camera. Please check again...')

# get the frame width and height
frame_width = int(cap.get(3))
frame_height = int(cap.get(4))

# define codec and create VideoWriter object
out = cv2.VideoWriter('../outputs/asl.mp4',
                      cv2.VideoWriter_fourcc(*'mp4v'), 30, (frame_width, frame_height))

# read until end of video
while(cap.isOpened()):
    # capture each frame of the video
    ret, frame = cap.read()
    # get the hand area on the video capture screen
    cv2.rectangle(frame, (100, 100), (324, 324), (20, 34, 255), 2)
    hand = hand_area(frame)
```

```
image = hand
```

```
image = np.transpose(image, (2, 0, 1)).astype(np.float32)
```

```
image = torch.tensor(image, dtype=torch.float)
```

```
# image = torch.tensor(image, dtype=torch.float).cuda()
```

```
image = image.unsqueeze(0)
```

```
outputs = model(image)
```

```
_, preds = torch.max(outputs.data, 1)
```

```
# print('PREDS', preds)
```

```
# print(f"Predicted output: {lb.classes_[preds]}")
```

```
cv2.putText(frame, lb.classes_[preds], (10, 30),
```

```
            cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0, 0, 255), 2)
```

```
cv2.imshow('image', frame)
```

```
out.write(frame)
```

```
# time.sleep(0.09)
```

```
# press 'q' to exit
```

```
if cv2.waitKey(27) & 0xFF == ord('q'):
```

```
    break
```

```
# release VideoCapture()
```

```
cap.release()
```

```
# close all frames and video windows
```

```
cv2.destroyAllWindows()
```

cnn_models.py

```
import torch.nn as nn
import torch.nn.functional as F
import joblib
```

```
from torchvision import models
```

```
# load the binarized labels
print('Loading label binarizer...')
lb = joblib.load('../outputs/lb.pkl')
```

```
class CustomCNN(nn.Module):
    def __init__(self):
        super(CustomCNN, self).__init__()
        self.conv1 = nn.Conv2d(3, 16, 5)
        self.conv2 = nn.Conv2d(16, 32, 5)
        self.conv3 = nn.Conv2d(32, 64, 3)
        self.conv4 = nn.Conv2d(64, 128, 5)

        self.fc1 = nn.Linear(128, 256)
        self.fc2 = nn.Linear(256, len(lb.classes_))

        self.pool = nn.MaxPool2d(2, 2)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = self.pool(F.relu(self.conv3(x)))
        x = self.pool(F.relu(self.conv4(x)))
        bs, _, _, _ = x.shape
        x = F.adaptive_avg_pool2d(x, 1).reshape(bs, -1)
        x = F.relu(self.fc1(x))
        x = self.fc2(x)
        return x
```

create_csv.py

```
import pandas as pd
import numpy as np
import os
import joblib

from sklearn.preprocessing import LabelBinarizer
from tqdm import tqdm
from imutils import paths

# get all the image paths
image_paths = list(paths.list_images('../input/preprocessed_image'))

# create a DataFrame
data = pd.DataFrame()

labels = []
for i, image_path in tqdm(enumerate(image_paths), total=len(image_paths)):
    label = image_path.split(os.path.sep)[-2]
    # save the relative path for mapping image to target
    data.loc[i, 'image_path'] = image_path

    labels.append(label)

labels = np.array(labels)
# one hot encode the labels
lb = LabelBinarizer()
labels = lb.fit_transform(labels)

print(f"The first one hot encoded labels: {labels[0]}")
print(f"Mapping the first one hot encoded label to its category: {lb.classes_[0]}")
print(f"Total instances: {len(labels)}")

for i in range(len(labels)):
    index = np.argmax(labels[i])
    data.loc[i, 'target'] = int(index)

# shuffle the dataset
data = data.sample(frac=1).reset_index(drop=True)

# save as CSV file
data.to_csv('../input/data.csv', index=False)

# pickle the binarized labels
print('Saving the binarized labels as pickled file')
joblib.dump(lb, '../outputs/lb.pkl')

print(data.head(10))
```

'''

USAGE:

```
python preprocess_image.py --num-images 1200
```

'''

```
import pandas as pd
import os
import cv2
import random
import albumentations
import numpy as np
import argparse

from imutils import paths
from tqdm import tqdm

parser = argparse.ArgumentParser()
parser.add_argument('-n', '--num-images', default=1000, type=int,
                    help='number of images to preprocess for each category')
args = vars(parser.parse_args())

print(f"Preprocessing {args['num_images']} from each category...")

# get all the image paths
image_paths = list(paths.list_images('../input/asl_alphabet_train/asl_alphabet_train'))
dir_paths = os.listdir('../input/asl_alphabet_train/asl_alphabet_train')
dir_paths.sort()

root_path = '../input/asl_alphabet_train/asl_alphabet_train'

# get 1000 images from each category
for idx, dir_path in tqdm(enumerate(dir_paths), total=len(dir_paths)):
    all_images = os.listdir(f"{root_path}/{dir_path}")
    os.makedirs(f"../input/preprocessed_image/{dir_path}", exist_ok=True)
    for i in range(args['num_images']): # how many images to preprocess for each category
        # generate a random id between 0 and 2999
        rand_id = (random.randint(0, 2999))
        image = cv2.imread(f"{root_path}/{dir_path}/{all_images[rand_id]}")
        image = cv2.resize(image, (224, 224))

        cv2.imwrite(f"../input/preprocessed_image/{dir_path}/{dir_path}{i}.jpg", image)

print('DONE')
```

"""

USAGE:

```
python test.py --img A_test.jpg
```

"""

```
import torch
import joblib
import torch.nn as nn
import numpy as np
import cv2
import argparse
import albumentations
import torch.nn.functional as F
import time
import cnn_models

# construct the argument parser and parse the arguments
parser = argparse.ArgumentParser()
parser.add_argument('-i', '--img', default='A_test.jpg', type=str,
                    help='path for the image to test on')
args = vars(parser.parse_args())

# load label binarizer
lb = joblib.load('../outputs/lb.pkl')

aug = albumentations.Compose([
    albumentations.Resize(224, 224, always_apply=True),
])

# model = models.MobNetV2(pretrained=False, requires_grad=False)
model = cnn_models.CustomCNN()
# model = cnn_models.CustomCNN().cuda()

model.load_state_dict(torch.load('../outputs/model.pth'))
print(model)
print('Model loaded')

image = cv2.imread(f"../input/asl_alphabet_test/{args['img']}")
image_copy = image.copy()

image = aug(image=np.array(image))['image']
image = np.transpose(image, (2, 0, 1)).astype(np.float32)
image = torch.tensor(image, dtype=torch.float)
# image = torch.tensor(image, dtype=torch.float).cuda()

image = image.unsqueeze(0)
print(image.shape)

start = time.time()
outputs = model(image)
_, preds = torch.max(outputs.data, 1)
print('PREDS', preds)
print(f"Predicted output: {lb.classes_[preds]}")
end = time.time()
print(f"{{(end-start):.3f}} seconds")
```

```
cv2.putText(image_copy, lb.classes_[preds], (10, 30), cv2.FONT_HERSHEY_SIMPLEX, 0.9, (0, 0, 255),
2)
cv2.imshow('image', image_copy)
cv2.imwrite(f"../outputs/{args['img']}", image_copy)
cv2.waitKey(0)
```

```
"""
```

USAGE:

```
python train.py --epochs 10
```

```
"""
```

```
import pandas as pd
import joblib
import numpy as np
import torch
import random
import albumentations
import matplotlib.pyplot as plt
import argparse
import torch.nn as nn
import torch.nn.functional as F
import torch.optim as optim
import torchvision.transforms as transforms
import time
import cv2
import cnn_models
```

```
from PIL import Image
from tqdm import tqdm
from sklearn.model_selection import train_test_split
from torch.utils.data import Dataset, DataLoader
```

```
# construct the argument parser and parse the arguments
parser = argparse.ArgumentParser()
parser.add_argument('-e', '--epochs', default=10, type=int,
                    help='number of epochs to train the model for')
args = vars(parser.parse_args())
```

```
""" SEED Everything """
def seed_everything(SEED=42):
    random.seed(SEED)
    np.random.seed(SEED)
    torch.manual_seed(SEED)
    torch.cuda.manual_seed(SEED)
    torch.cuda.manual_seed_all(SEED)
    torch.backends.cudnn.benchmark = True
SEED=42
seed_everything(SEED=SEED)
""" SEED Everything """
```

```
# set computation device
device = ('cuda:0' if torch.cuda.is_available() else 'cpu')
# device =(torch.cuda.get_device_properties(0) if torch.cuda.is_available() else 'CPU')
```

```
print(f"Computation device: {device}")
```

```
# read the data.csv file and get the image paths and labels
df = pd.read_csv('../input/data.csv')
X = df.image_path.values
```



```
y = df.target.values
```

```
(xtrain, xtest, ytrain, ytest) = (train_test_split(X, y,  
                                                test_size=0.15, random_state=42))
```

```
print(f"Training on {len(xtrain)} images")  
print(f"Validation on {len(xtest)} images")
```

```
# image dataset module
```

```
class ASLImageDataset(Dataset):
```

```
    def __init__(self, path, labels):
```

```
        self.X = path
```

```
        self.y = labels
```

```
        # apply augmentations
```

```
        self.aug = albumentations.Compose([  
            albumentations.Resize(224, 224, always_apply=True),  
        ])
```

```
    def __len__(self):
```

```
        return (len(self.X))
```

```
    def __getitem__(self, i):
```

```
        # image = Image.open(self.X[i])
```

```
        image = cv2.imread(self.X[i])
```

```
        image = self.aug(image=np.array(image))['image']
```

```
        image = np.transpose(image, (2, 0, 1)).astype(np.float32)
```

```
        label = self.y[i]
```

```
        return torch.tensor(image, dtype=torch.float), torch.tensor(label, dtype=torch.long)
```

```
train_data = ASLImageDataset(xtrain, ytrain)
```

```
test_data = ASLImageDataset(xtest, ytest)
```

```
# dataloaders
```

```
trainloader = DataLoader(train_data, batch_size=32, shuffle=True)
```

```
testloader = DataLoader(test_data, batch_size=32, shuffle=False)
```

```
# model = models.MobileNetV2(pretrained=True, requires_grad=False)
```

```
model = cnn_models.CustomCNN().to(device)
```

```
print(model)
```

```
# Find total parameters and trainable parameters
```

```
total_params = sum(p.numel() for p in model.parameters())
```

```
print(f"{total_params:,} total parameters.")
```

```
total_trainable_params = sum(
```

```
    p.numel() for p in model.parameters() if p.requires_grad)
```

```
print(f"{total_trainable_params:,} training parameters.")
```

```
# optimizer
```

```
optimizer = optim.Adam(model.parameters(), lr=0.001)
```

```
# loss function
```

```
criterion = nn.CrossEntropyLoss()
```

```
#validation function
```

```

def validate(model, dataloader):
    print('Validating')
    model.eval()
    running_loss = 0.0
    running_correct = 0
    with torch.no_grad():
        for i, data in tqdm(enumerate(dataloader), total=int(len(test_data)/dataloader.batch_size)):
            data, target = data[0].to(device), data[1].to(device)
            outputs = model(data)
            loss = criterion(outputs, target)

            running_loss += loss.item()
            _, preds = torch.max(outputs.data, 1)
            running_correct += (preds == target).sum().item()

    val_loss = running_loss/len(dataloader.dataset)
    val_accuracy = 100. * running_correct/len(dataloader.dataset)
    print(f'Val Loss: {val_loss:.4f}, Val Acc: {val_accuracy:.2f}')

    return val_loss, val_accuracy

```

```

# training function
def fit(model, dataloader):
    print('Training')
    model.train()
    running_loss = 0.0
    running_correct = 0
    for i, data in tqdm(enumerate(dataloader), total=int(len(train_data)/dataloader.batch_size)):
        data, target = data[0].to(device), data[1].to(device)
        optimizer.zero_grad()
        outputs = model(data)
        loss = criterion(outputs, target)
        running_loss += loss.item()
        _, preds = torch.max(outputs.data, 1)
        # print(preds)
        running_correct += (preds == target).sum().item()
        loss.backward()
        optimizer.step()

    train_loss = running_loss/len(dataloader.dataset)
    train_accuracy = 100. * running_correct/len(dataloader.dataset)

    print(f"Train Loss: {train_loss:.4f}, Train Acc: {train_accuracy:.2f}")

    return train_loss, train_accuracy

```

```

train_loss , train_accuracy = [], []
val_loss , val_accuracy = [], []
start = time.time()
for epoch in range(args['epochs']):
    print(f"Epoch {epoch+1} of {args['epochs']}")
    train_epoch_loss, train_epoch_accuracy = fit(model, trainloader)
    val_epoch_loss, val_epoch_accuracy = validate(model, testloader)
    train_loss.append(train_epoch_loss)
    train_accuracy.append(train_epoch_accuracy)

```

```
val_loss.append(val_epoch_loss)
val_accuracy.append(val_epoch_accuracy)
end = time.time()
```

```
print(f"{{(end-start)/60:.3f}} minutes")
```

```
# accuracy plots
```

```
plt.figure(figsize=(10, 7))
plt.plot(train_accuracy, color='green', label='train accuracy')
plt.plot(val_accuracy, color='blue', label='validataion accuracy')
plt.xlabel('Epochs')
plt.ylabel('Accuracy')
plt.legend()
plt.savefig('../outputs/accuracy.png')
plt.show()
```

```
# loss plots
```

```
plt.figure(figsize=(10, 7))
plt.plot(train_loss, color='orange', label='train loss')
plt.plot(val_loss, color='red', label='validataion loss')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend()
plt.savefig('../outputs/loss.png')
plt.show()
```

```
# save the model to disk
```

```
print('Saving model...')
torch.save(model.state_dict(), '../outputs/model.pth')
```